

Architecture: End-to-End Fraud Detection System

This document outlines the complete architecture of the automated fraud detection system. The system is designed not just to build a model, but to *discover the optimal model* through intelligent hyperparameter tuning and then deploy it for real-time use.

The system is broken into two primary phases:

1. **Phase 1: The "Search & Build" Phase** (The Training Pipeline)
2. **Phase 2: The "Deploy & Predict" Phase** (The Inference App)

Phase 1: The "Search & Build" Phase (Training)

This phase is handled by the `fraud_pipeline_complete.py` script. Its goal is to run an exhaustive search to find the best possible combination of pipeline steps and model parameters, then build and save the single best-performing model.

The entire process is managed by **Optuna**, a hyperparameter optimization framework.

The Optuna Tuning Loop

The script executes 50 "trials" (or more). In each trial, Optuna intelligently suggests a new set of parameters and builds an *entire* ML pipeline from scratch to test them.

Here is the architecture of a **single trial**:

1. **Optuna Suggests Parameters:**
 - `n_features_to_select` : How many features (10-25) should RFE keep?
 - `smote_k_neighbors` : How many neighbors (3-7) should SMOTE use?
 - `knn_n_neighbors` : How many neighbors (3-15) for the KNN model?
 - `rf_n_estimators` : How many trees (100-300) for the Random Forest?
 - `rf_max_depth` : How deep (10-30) can the trees be?
 - `meta_C` : How strong (0.01-10.0) should the meta-model's regularization be?
2. **A Temporary Pipeline is Built:** This pipeline defines the *order of operations*, which is critical for preventing data leakage.

[Input Data (`X_train`)] → [Step 1] → [Step 2] → [Step 3] → [Step 4] → [Score]

- **Step 1:** RobustScaler
 - Scales the `Time` and `Amount` features, making them robust to outliers.
- **Step 2: RFE (Recursive Feature Elimination)**
 - Uses a simple `LogisticRegression` model to analyze all 30 features.
 - It recursively removes the weakest features until only the `n_features_to_select` (e.g., 13) best ones remain.
- **Step 3: SMOTE (Synthetic Minority Over-sampling)**
 - Analyzes the rare "Fraud" cases.

- Creates new, synthetic "Fraud" data points based on its neighbors (`smote_k_neighbors`). This balances the dataset so the model can learn what fraud looks like.
- **Step 4: StackingClassifier (The "Brain" of the Model)**
 - This is a 2-level ensemble model that runs its own internal parallel processes (`n_jobs=-1`) to be fast.
 - See the **Stacking Architecture** diagram below.

3. Evaluation:

- This entire pipeline is evaluated using 3-fold cross-validation on the training data.
- The outer loop (`cross_val_score`) runs in serial (`n_jobs=1`), but the inner `StackingClassifier` runs in parallel (`n_jobs=-1`). This is the "parallelism fix" that allows the CPU to hit 100% in bursts for maximum efficiency.
- The average `roc_auc` score (a metric for imbalance) is returned to Optuna.

4. Learn & Repeat:

- Optuna records the result and uses it to make a smarter guess for the next trial.
- This loop repeats 50 times.

End of Phase 1: Building the Final Model

After the loop, Optuna knows the single best set of parameters (e.g., the 13 features, 7 KNN neighbors, etc. from your output).

1. **Build Final Pipeline:** The script builds *one last pipeline* using these winning parameters.
2. **Train:** It trains this final pipeline on the *entire* training dataset.
3. **Save:** The fully trained `final_pipeline` object is saved to `fraud_detection_pipeline.joblib`, and the list of feature names is saved to `feature_names.joblib`.

The Stacking Architecture (The "Brain")

This is the detailed view of "Step 4" from the pipeline. The `StackingClassifier` itself is an ensemble of models.

1. Level 0: The "Workers" (Base Models)

- The same data (after SMOTE) is fed to three different models in parallel.
- `KNeighborsClassifier` (KNN)
- `RandomForestClassifier` (RF)
- `AdaBoostClassifier` (AdaBoost)
- These models don't give the final answer. They just give their own "opinion" (prediction).

2. Level 1: The "Manager" (Meta-Model)

- A simple `LogisticRegression` model.
- Its only job is to look at the predictions from the three "workers."
- It *learns* which workers to trust. For example, it might learn "When RF and AdaBoost agree, I should trust them. When KNN disagrees, I should ignore it."

- It makes the final decision, which is more robust than any single model.

Phase 2: The "Deploy & Predict" Phase (Inference)

This phase is handled by the `app.py` (Streamlit) file. It uses the saved files from Phase 1 to make live predictions.

1. On Startup:

- The Streamlit app loads `fraud_detection_pipeline.joblib` and `feature_names.joblib` into memory.

2. User Input:

- A user enters transaction data (`Amount` , `V4` , `V10` , etc.) into the web-page sliders and input boxes.

3. Real-Time Prediction:

- The app creates a single-row `DataFrame` with the user's data, using `feature_names.joblib` to ensure the column order is correct.
- This `DataFrame` is fed into the loaded `final_pipeline.predict_proba()` method.

4. The "Inference" Pipeline:

- The loaded pipeline automatically runs all its saved steps on the new data:
- [New Data] -> [Step 1] -> [Step 2] -> [Step 3] -> [Prediction]
- **Step 1:** RobustScaler
 - Transforms the new `Time` and `Amount` using the same scaling it learned from the training data.
- **Step 2:** RFE
 - Instantly selects only the 13 best features it learned to pick during training.
- **Step 3:** SMOTE
 - **This step is intelligently SKIPPED.** SMOTE is only for *training*, not for predicting. The pipeline knows this automatically.
- **Step 4:** StackingClassifier
 - The new data goes to the "Workers" (KNN, RF, AdaBoost), who all make a prediction.
 - Their predictions go to the "Manager" (LogisticRegression), which makes the final decision.

5. Display Result:

- The app gets the final fraud probability (e.g., 87%) and displays either a "FRAUD DETECTED" error or a "Normal Transaction" success message to the user in their browser.